

Different architectures for CNN on MNIST dataset

1. Import Libraries

In [1]:

```
1 # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py
2 from __future__ import print_function
3 import keras
4 from keras.datasets import mnist
5 from keras.models import Sequential
6 from keras.layers import Dense, Dropout, Flatten, BatchNormalization, Conv2D, MaxPooling2D
7 from keras import layers
8 from keras import backend as K
9 import matplotlib.pyplot as plt
10 import pandas as pd
11 %matplotlib inline
```

Using TensorFlow backend.

2. Global variables

In [39]:

```
1 batch_size = 128
2 num_classes = 10
3 epochs = 10
4 img_rows, img_cols, channels = 28, 28, 1
5 model_results = dict()
```

3. Load Data

In [3]:

```
1 # the data, split between train and test sets
2 (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

3.1. Normalising Data

In [4]:

```
1 if K.image_data_format() == 'channels_first':
2     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
3     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
4     input_shape = (channels, img_rows, img_cols)
5 else:
6     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
7     x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
8     input_shape = (img_rows, img_cols, channels)
9
10 x_train = x_train.astype('float32')
11 x_test = x_test.astype('float32')
12 x_train /= 255
13 x_test /= 255
14 print('x_train shape:', x_train.shape)
15 print(x_train.shape[0], 'train samples')
16 print(x_test.shape[0], 'test samples')
17
18 # convert class vectors to binary class matrices
19 y_train = keras.utils.to_categorical(y_train, num_classes)
20 y_test = keras.utils.to_categorical(y_test, num_classes)
```

x_train shape: (60000, 28, 28, 1)

60000 train samples

10000 test samples

4.1. Model 1 - 2 Conv2D layers, 1 Max pool layer, 2 dense layers + dropout and flatten

In [5]:

```
1 model = Sequential()
2
3 #IN: (batch,rows,columns,channel)=>(batch,28,28,1)
4 #Remember the input_shape is specified without the batch
5 model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
6 #OUT: (batch,rows,columns,channel)=>(batch,26,26,32)
7 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
8 #OUT: (batch,rows,columns,channel)=>(batch,24,24,64)
9 model.add(MaxPooling2D(pool_size=(2, 2)))
10 #OUT: (batch,rows,columns,channel)=>(batch,12,12,64)
11 model.add(Dropout(0.25))
12 model.add(Flatten())
13 #OUT: (batch,rows,columns,channel)=>(batch,9216)
14 model.add(Dense(128, activation='relu'))
15 model.add(Dropout(0.5))
16 model.add(Dense(num_classes, activation='softmax'))
```

WARNING:tensorflow:From c:\users\byron\applications\pythonmaster\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From c:\users\byron\applications\pythonmaster\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

In [6]:

```
1 model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
dense_1 (Dense)	(None, 128)	1179776
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
Total params: 1,199,882		
Trainable params: 1,199,882		
Non-trainable params: 0		

In [7]:

```
1 # model.get_config()
2 # model.inputs
3 # model.outputs
4 # model.get_weights()
```

In [8]:

```
1 model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.A
```

In [9]:

```
1 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1,
```

WARNING:tensorflow:From c:\users\byron\applications\pythonmaster\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 157s 3ms/step - loss: 0.2373
- acc: 0.9268 - val_loss: 0.0497 - val_acc: 0.9841

Epoch 2/10

60000/60000 [=====] - 156s 3ms/step - loss: 0.0858
- acc: 0.9747 - val_loss: 0.0349 - val_acc: 0.9874

Epoch 3/10

60000/60000 [=====] - 160s 3ms/step - loss: 0.0635
- acc: 0.9810 - val_loss: 0.0344 - val_acc: 0.9893

Epoch 4/10

60000/60000 [=====] - 157s 3ms/step - loss: 0.0519
- acc: 0.9840 - val_loss: 0.0284 - val_acc: 0.9903

Epoch 5/10

60000/60000 [=====] - 158s 3ms/step - loss: 0.0442
- acc: 0.9862 - val_loss: 0.0277 - val_acc: 0.9908

Epoch 6/10

60000/60000 [=====] - 87s 1ms/step - loss: 0.0386 -
acc: 0.9878 - val_loss: 0.0302 - val_acc: 0.9902

Epoch 7/10

60000/60000 [=====] - 84s 1ms/step - loss: 0.0332 -
acc: 0.9898 - val_loss: 0.0267 - val_acc: 0.9906

Epoch 8/10

60000/60000 [=====] - 85s 1ms/step - loss: 0.0322 -
acc: 0.9901 - val_loss: 0.0275 - val_acc: 0.9909

Epoch 9/10

60000/60000 [=====] - 86s 1ms/step - loss: 0.0286 -
acc: 0.9907 - val_loss: 0.0248 - val_acc: 0.9928

Epoch 10/10


60000/60000 [=====] - 86s 1ms/step - loss: 0.0255 -
acc: 0.9920 - val_loss: 0.0246 - val_acc: 0.9931

In [10]:

```

1 x = list(range(1,epochs+1))
2 plt.plot(x, history.history['loss'],label='Train Loss')
3 plt.plot(x, history.history['val_loss'],label='CV Loss')
4 plt.xlabel('epoch')
5 plt.ylabel('Categorical Crossentropy Loss')
6 plt.legend()
7 plt.show()

```



In [11]:

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test loss:', score[0])
3 print('Test accuracy:', score[1])

```

Test loss: 0.02459182265629006

Test accuracy: 0.9931

In [12]:

```

1 score = model.evaluate(x_train, y_train, verbose=0)
2 print('Train loss:', score[0])
3 print('Train accuracy:', score[1])

```

Train loss: 0.004404588298801052

Train accuracy: 0.9988833333333333

In [40]:

```

1 model_results['model1'] = {'Test Loss': 0.0245, 'Train Loss':0.0044}

```

4.2. Model 2 - 2 Conv2D layers, 2 Max pool layers, 2 dense layers

+ dropout, flatten and padding

In [14]:

```

1 model = Sequential()
2
3 model.add(Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu',
4                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
5
6 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
7
8 model.add(Conv2D(filters=32, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu',
9                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
10
11 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
12
13 model.add(Flatten())
14
15 model.add(Dense(units=100, activation=keras.activations.relu, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
16
17 model.add(Dropout(0.5))
18
19 model.add(Dense(units=num_classes, activation=keras.activations.softmax, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))

```

In [15]:

```
1 model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_4 (Conv2D)	(None, 14, 14, 32)	9248
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 32)	0
flatten_2 (Flatten)	(None, 1568)	0
dense_3 (Dense)	(None, 100)	156900
dropout_3 (Dropout)	(None, 100)	0
dense_4 (Dense)	(None, 10)	1010
=====		
Total params: 167,478		
Trainable params: 167,478		
Non-trainable params: 0		

In [16]:

```
1 model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.A
```

In [17]:

```
1 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1,
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 34s 567us/step - loss: 0.2788

- acc: 0.9140 - val_loss: 0.0575 - val_acc: 0.9817

Epoch 2/10

60000/60000 [=====] - 33s 557us/step - loss: 0.0984

- acc: 0.9710 - val_loss: 0.0450 - val_acc: 0.9845

Epoch 3/10

60000/60000 [=====] - 34s 569us/step - loss: 0.0745

- acc: 0.9779 - val_loss: 0.0380 - val_acc: 0.9877

Epoch 4/10

60000/60000 [=====] - 34s 559us/step - loss: 0.0602

- acc: 0.9822 - val_loss: 0.0314 - val_acc: 0.9887

Epoch 5/10

60000/60000 [=====] - 34s 559us/step - loss: 0.0540

- acc: 0.9839 - val_loss: 0.0316 - val_acc: 0.9905

Epoch 6/10

60000/60000 [=====] - 34s 559us/step - loss: 0.0462

- acc: 0.9860 - val_loss: 0.0320 - val_acc: 0.9896

Epoch 7/10

60000/60000 [=====] - 34s 563us/step - loss: 0.0410

- acc: 0.9872 - val_loss: 0.0282 - val_acc: 0.9910

Epoch 8/10

60000/60000 [=====] - 34s 564us/step - loss: 0.0384

- acc: 0.9877 - val_loss: 0.0314 - val_acc: 0.9905

Epoch 9/10

60000/60000 [=====] - 34s 562us/step - loss: 0.0344

- acc: 0.9892 - val_loss: 0.0282 - val_acc: 0.9898

Epoch 10/10

60000/60000 [=====] - 34s 572us/step - loss: 0.0335


- acc: 0.9895 - val_loss: 0.0261 - val_acc: 0.9910

In [18]:

```

1 x = list(range(1,epochs+1))
2 plt.plot(x, history.history['loss'],label='Train Loss')
3 plt.plot(x, history.history['val_loss'],label='CV Loss')
4 plt.xlabel('epoch')
5 plt.ylabel('Categorical Crossentropy Loss')
6 plt.legend()
7 plt.show()

```



In [19]:

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test loss:', score[0])
3 print('Test accuracy:', score[1])

```

Test loss: 0.026057142214136592

Test accuracy: 0.991

In [20]:

```

1 score = model.evaluate(x_train, y_train, verbose=0)
2 print('Train loss:', score[0])
3 print('Train accuracy:', score[1])

```

Train loss: 0.009017451262661794

Train accuracy: 0.9971333333333333

In [41]:

```

1 model_results['model2'] = {'Test Loss': 0.0260, 'Train Loss': 0.0090}

```

4.3. Model 3 - 4 Conv2D layers, 2 Max pool layers, 2 dense layers

+ dropout,flatten,padding and batchnormalising

In [22]:

```
1 model = Sequential()
2
3 #=====CONV BLOCK 1=====
4 model.add(Conv2D(filters=64, kernel_size=(2,2), strides=(1,1), padding='same', activation='relu',
5                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
6 model.add(BatchNormalization())
7 model.add(Conv2D(filters=64, kernel_size=(2,2), strides=(1,1), padding='valid', activation='relu',
8                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
9 model.add(BatchNormalization())
10 model.add(MaxPooling2D(pool_size=(3,3), strides=(3,3)))
11 #=====CONV BLOCK 1=====
12
13
14 #=====CONV BLOCK 2=====
15 model.add(Conv2D(filters=128, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu',
16                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
17 model.add(BatchNormalization())
18 model.add(Conv2D(filters=128, kernel_size=(3,3), strides=(1,1), padding='same', activation='relu',
19                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
20 model.add(BatchNormalization())
21 model.add(MaxPooling2D(pool_size=(3,3), strides=(3,3)))
22 #=====CONV BLOCK 2=====
23
24
25 #=====DENSE=====
26 model.add(Flatten())
27 model.add(Dense(units=200, activation=keras.activations.relu, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
28 model.add(Dropout(0.5))
29 model.add(BatchNormalization())
30 model.add(Dense(units=num_classes, activation=keras.activations.softmax, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer='zeros'))
31 #=====DENSE=====
```

In [23]:

```
1 model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 28, 28, 64)	320
batch_normalization_1 (Batch Normalization)	(None, 28, 28, 64)	256
conv2d_6 (Conv2D)	(None, 27, 27, 64)	16448
batch_normalization_2 (Batch Normalization)	(None, 27, 27, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 9, 9, 64)	0
conv2d_7 (Conv2D)	(None, 9, 9, 128)	73856
batch_normalization_3 (Batch Normalization)	(None, 9, 9, 128)	512
conv2d_8 (Conv2D)	(None, 9, 9, 128)	147584
batch_normalization_4 (Batch Normalization)	(None, 9, 9, 128)	512
max_pooling2d_5 (MaxPooling2D)	(None, 3, 3, 128)	0
flatten_3 (Flatten)	(None, 1152)	0
dense_5 (Dense)	(None, 200)	230600
dropout_4 (Dropout)	(None, 200)	0
batch_normalization_5 (Batch Normalization)	(None, 200)	800
dense_6 (Dense)	(None, 10)	2010
=====		
Total params: 473,154		
Trainable params: 471,986		
Non-trainable params: 1,168		

In [24]:

```
1 model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.A
```

In [25]:

```
1 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1,
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 554s 9ms/step - loss: 0.1633

- acc: 0.9551 - val_loss: 0.0503 - val_acc: 0.9844

Epoch 2/10

60000/60000 [=====] - 552s 9ms/step - loss: 0.0473

- acc: 0.9877 - val_loss: 0.0340 - val_acc: 0.9884

Epoch 3/10

60000/60000 [=====] - 551s 9ms/step - loss: 0.0350

- acc: 0.9904 - val_loss: 0.0270 - val_acc: 0.9907

Epoch 4/10

60000/60000 [=====] - 553s 9ms/step - loss: 0.0271

- acc: 0.9920 - val_loss: 0.0306 - val_acc: 0.9906

Epoch 5/10

60000/60000 [=====] - 552s 9ms/step - loss: 0.0221

- acc: 0.9939 - val_loss: 0.0200 - val_acc: 0.9935

Epoch 6/10

60000/60000 [=====] - 552s 9ms/step - loss: 0.0182

- acc: 0.9947 - val_loss: 0.0318 - val_acc: 0.9898

Epoch 7/10

60000/60000 [=====] - 552s 9ms/step - loss: 0.0162

- acc: 0.9950 - val_loss: 0.0219 - val_acc: 0.9934

Epoch 8/10

60000/60000 [=====] - 556s 9ms/step - loss: 0.0156

- acc: 0.9952 - val_loss: 0.0239 - val_acc: 0.9921

Epoch 9/10

60000/60000 [=====] - 552s 9ms/step - loss: 0.0131

- acc: 0.9960 - val_loss: 0.0241 - val_acc: 0.9917

Epoch 10/10

60000/60000 [=====] - 553s 9ms/step - loss: 0.0114


- acc: 0.9966 - val_loss: 0.0184 - val_acc: 0.9944

In [26]:

```

1 x = list(range(1,epochs+1))
2 plt.plot(x, history.history['loss'],label='Train Loss')
3 plt.plot(x, history.history['val_loss'],label='CV Loss')
4 plt.xlabel('epoch')
5 plt.ylabel('Categorical Crossentropy Loss')
6 plt.legend()
7 plt.show()

```



In [27]:

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test loss:', score[0])
3 print('Test accuracy:', score[1])

```

Test loss: 0.01836347421152368

Test accuracy: 0.9944

In [28]:

```

1 score = model.evaluate(x_train, y_train, verbose=0)
2 print('Train loss:', score[0])
3 print('Train accuracy:', score[1])

```

Train loss: 0.003370044193206043

Train accuracy: 0.9988666666666667

In [42]:

```

1 model_results['model3'] = {'Test Loss': 0.0183, 'Train Loss':0.0033}

```

I think the above model overfits somewhat due to the deep network

4.4. Model 4 - 2 Conv2D layers, 1 Max pool layers, 2 dense layers + dropout,flatten,padding,batchnormalising and stride

In [30]:

```
1 model = Sequential()
2
3 #=====CONV BLOCK 1=====
4 model.add(Conv2D(filters=64, kernel_size=(2,2), strides=(2,2), padding='same', activation='relu',
5                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer = keras.initializers.zeros))
6 model.add(BatchNormalization())
7 model.add(Conv2D(filters=64, kernel_size=(4,4), strides=(2,2), padding='same', activation='relu',
8                 kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer = keras.initializers.zeros))
9 model.add(BatchNormalization())
10 model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2)))
11 #=====CONV BLOCK 1=====
12
13
14 #=====DENSE=====
15 model.add(Flatten())
16 model.add(Dense(units=200, activation=keras.activations.relu, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer = keras.initializers.zeros))
17 model.add(Dropout(0.5))
18 model.add(BatchNormalization())
19 model.add(Dense(units=num_classes, activation=keras.activations.softmax, kernel_initializer = keras.initializers.he_uniform(seed=1), bias_initializer = keras.initializers.zeros))
20 #=====DENSE=====
```

In [31]:

```
1 model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_9 (Conv2D)	(None, 14, 14, 64)	320
batch_normalization_6 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_10 (Conv2D)	(None, 7, 7, 64)	65600
batch_normalization_7 (Batch Normalization)	(None, 7, 7, 64)	256
max_pooling2d_6 (MaxPooling2D)	(None, 3, 3, 64)	0
flatten_4 (Flatten)	(None, 576)	0
dense_7 (Dense)	(None, 200)	115400
dropout_5 (Dropout)	(None, 200)	0
batch_normalization_8 (Batch Normalization)	(None, 200)	800
dense_8 (Dense)	(None, 10)	2010
=====		
Total params: 184,642		
Trainable params: 183,986		
Non-trainable params: 656		
=====		

In [32]:

```
1 model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.A
```

In [33]:

```
1 history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1,
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.2895 -

acc: 0.9128 - val_loss: 0.0791 - val_acc: 0.9749

Epoch 2/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0971 -

acc: 0.9712 - val_loss: 0.0516 - val_acc: 0.9840

Epoch 3/10

60000/60000 [=====] - 72s 1ms/step - loss: 0.0696 -

acc: 0.9787 - val_loss: 0.0451 - val_acc: 0.9852

Epoch 4/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0589 -

acc: 0.9821 - val_loss: 0.0454 - val_acc: 0.9838

Epoch 5/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0474 -

acc: 0.9857 - val_loss: 0.0415 - val_acc: 0.9865

Epoch 6/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0415 -

acc: 0.9871 - val_loss: 0.0426 - val_acc: 0.9860

Epoch 7/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0353 -

acc: 0.9886 - val_loss: 0.0386 - val_acc: 0.9880

Epoch 8/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0319 -

acc: 0.9901 - val_loss: 0.0365 - val_acc: 0.9887

Epoch 9/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0283 -

acc: 0.9911 - val_loss: 0.0423 - val_acc: 0.9880

Epoch 10/10

60000/60000 [=====] - 73s 1ms/step - loss: 0.0256 -


acc: 0.9914 - val_loss: 0.0373 - val_acc: 0.9890

In [34]:

```

1 x = list(range(1,epochs+1))
2 plt.plot(x, history.history['loss'],label='Train Loss')
3 plt.plot(x, history.history['val_loss'],label='CV Loss')
4 plt.xlabel('epoch')
5 plt.ylabel('Categorical Crossentropy Loss')
6 plt.legend()
7 plt.show()

```



In [35]:

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test loss:', score[0])
3 print('Test accuracy:', score[1])

```

Test loss: 0.03733194876129637

Test accuracy: 0.989

In [36]:

```

1 score = model.evaluate(x_train, y_train, verbose=0)
2 print('Train loss:', score[0])
3 print('Train accuracy:', score[1])

```

Train loss: 0.007702990999433435

Train accuracy: 0.99775

In [45]:

```

1 model_results['model4'] = {'Test Loss': 0.0373, 'Train Loss': 0.0077}

```

5. Model Results

In [46]:

```
1 pd.DataFrame(model_results)
```

Out[46]:

	model1	model2	model3	model4
Test Loss	0.0245	0.026	0.0183	0.0373
Train Loss	0.0044	0.009	0.0033	0.0077