

1. Keras - MLPs on MNIST

2. Import Libraries

In [1]:

```
1 # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
2 from keras.utils import np_utils
3 from keras.datasets import mnist
4 from keras.optimizers import Adam,RMSprop,SGD
5 from keras.layers import Dropout
6 from keras.layers import Dense, Activation
7 from keras.layers.normalization import BatchNormalization
8 from keras.models import Sequential
9 from keras.initializers import RandomNormal
10 from keras.wrappers.scikit_learn import KerasClassifier
11 from sklearn.model_selection import GridSearchCV
12 import seaborn as sns
13 import matplotlib.pyplot as plt
14 import numpy as np
15 import time
```

Using TensorFlow backend.

3. Plot Function

In [2]:

```
1 # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
2 # https://stackoverflow.com/a/14434334
3 # this function is used to update the plots for each epoch and error
4 def plt_dynamic(x, vy, ty, ax, colors=['b']):
5     ax.plot(x, vy, 'b', label="Validation Loss")
6     ax.plot(x, ty, 'r', label="Train Loss")
7     plt.legend()
8     plt.grid()
9     fig.canvas.draw()
10
11
```

4. Data

In [3]:

```
1  # the data, shuffled and split between train and test sets
2  (X_train, y_train), (X_test, y_test) = mnist.load_data()
3
4  print("Number of training examples :", X_train.shape[0], "and each image is of shape (",
5  print("Number of training examples :", X_test.shape[0], "and each image is of shape (",
6
7  # if you observe the input shape its 2 dimensional vector
8  # for each image we have a (28*28) vector
9  # we will convert the (28*28) vector into single dimensional vector of 1 * 784
10
11 X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
12 X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
13
14 # after converting the input images from 3d to 2d vectors
15
16 print("Number of training examples :", X_train.shape[0], "and each image is of shape (",
17 print("Number of training examples :", X_test.shape[0], "and each image is of shape (",
18
19
```

```
Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)
Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)
```

In [4]:

```
1 # An example data point
2 print(X_train[0]);
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
```

5. Normalise the data

7.1 Parameters

In [7]:

```
1 output_dim = 10
2 input_dim = X_train.shape[1]
3
4 batch_size = 128
5 nb_epoch = 10
6
7
```

7.2.1 MLP (2 hidden layers) + Batch-Norm + Dropout + AdamOptimizer

In [8]:

```

1
2 model_1 = Sequential()
3
4 model_1.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=
5 model_1.add(BatchNormalization())
6 model_1.add(Dropout(0.5))
7
8 model_1.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
9 model_1.add(BatchNormalization())
10 model_1.add(Dropout(0.5))
11
12 model_1.add(Dense(output_dim, activation='softmax'))
13
14 model_1.summary()

```

WARNING:tensorflow:From C:\Users\Byron\AppData\Local\PythonMaster\lib\site-packages\tensorflow\python\framework\op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From C:\Users\Byron\AppData\Local\PythonMaster\lib\site-packages\keras\backend\tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
=====		
dense_1 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290
=====		
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [9]:

```

1 model_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

WARNING:tensorflow:From C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 8s 136us/step - loss: 0.4271
- acc: 0.8700 - val_loss: 0.1373 - val_acc: 0.9563

Epoch 2/10

60000/60000 [=====] - 7s 124us/step - loss: 0.2069
- acc: 0.9383 - val_loss: 0.1020 - val_acc: 0.9682

Epoch 3/10

60000/60000 [=====] - 7s 121us/step - loss: 0.1579
- acc: 0.9527 - val_loss: 0.0952 - val_acc: 0.9699

Epoch 4/10

60000/60000 [=====] - 7s 120us/step - loss: 0.1390
- acc: 0.9589 - val_loss: 0.0831 - val_acc: 0.9735

Epoch 5/10

60000/60000 [=====] - 7s 121us/step - loss: 0.1212
- acc: 0.9623 - val_loss: 0.0769 - val_acc: 0.9758

Epoch 6/10

60000/60000 [=====] - 7s 121us/step - loss: 0.1085
- acc: 0.9658 - val_loss: 0.0799 - val_acc: 0.9760

Epoch 7/10

60000/60000 [=====] - 7s 122us/step - loss: 0.1002
- acc: 0.9689 - val_loss: 0.0713 - val_acc: 0.9779

Epoch 8/10

60000/60000 [=====] - 7s 122us/step - loss: 0.0941
- acc: 0.9710 - val_loss: 0.0641 - val_acc: 0.9804

Epoch 9/10

60000/60000 [=====] - 7s 122us/step - loss: 0.0868
- acc: 0.9737 - val_loss: 0.0645 - val_acc: 0.9804

Epoch 10/10

60000/60000 [=====] - 7s 124us/step - loss: 0.0802
- acc: 0.9746 - val_loss: 0.0614 - val_acc: 0.9819

In [10]:


```

1 score = model_1.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.061443126168311574

Test accuracy: 0.9819



In [11]:

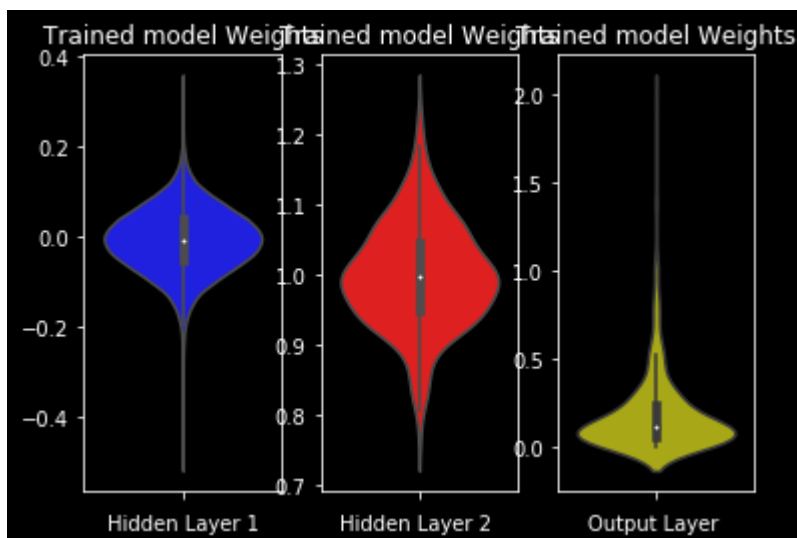
```

1  w_after = model_1.get_weights()
2
3  h1_w = w_after[0].flatten().reshape(-1,1)
4  h2_w = w_after[2].flatten().reshape(-1,1)
5  out_w = w_after[4].flatten().reshape(-1,1)
6
7
8  fig = plt.figure()
9  plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()
25
26 del(model_1)

```

C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\scipy\stats\stat
s.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional in
dexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the fu
ture this will be interpreted as an array index, `arr[np.array(seq)]`, which
will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



7.2.2 MLP (3 hidden layers) + Batch-Norm + Dropout + AdamOptimizer

In [12]:

```

1
2 model_2 = Sequential()
3
4 model_2.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
5 model_2.add(BatchNormalization())
6 model_2.add(Dropout(0.5))
7
8 model_2.add(Dense(300, activation='relu', kernel_initializer='he_normal'))
9 model_2.add(BatchNormalization())
10 model_2.add(Dropout(0.5))
11
12 model_2.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
13 model_2.add(BatchNormalization())
14 model_2.add(Dropout(0.5))
15
16 model_2.add(Dense(output_dim, activation='softmax'))
17
18 model_2.summary()
19

```

Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 500)	392500
batch_normalization_3 (Batch Normalization)	(None, 500)	2000
dropout_3 (Dropout)	(None, 500)	0
dense_5 (Dense)	(None, 300)	150300
batch_normalization_4 (Batch Normalization)	(None, 300)	1200
dropout_4 (Dropout)	(None, 300)	0
dense_6 (Dense)	(None, 100)	30100
batch_normalization_5 (Batch Normalization)	(None, 100)	400
dropout_5 (Dropout)	(None, 100)	0
dense_7 (Dense)	(None, 10)	1010
=====		
Total params: 577,510		
Trainable params: 575,710		
Non-trainable params: 1,800		

In [13]:

```

1 model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 10s 162us/step - loss: 0.5704
 - acc: 0.8269 - val_loss: 0.1611 - val_acc: 0.9505

Epoch 2/10

60000/60000 [=====] - 9s 146us/step - loss: 0.2497
 - acc: 0.9281 - val_loss: 0.1158 - val_acc: 0.9645

Epoch 3/10

60000/60000 [=====] - 9s 147us/step - loss: 0.1918
 - acc: 0.9439 - val_loss: 0.1005 - val_acc: 0.9705

Epoch 4/10

60000/60000 [=====] - 9s 145us/step - loss: 0.1587
 - acc: 0.9531 - val_loss: 0.0930 - val_acc: 0.9725

Epoch 5/10

60000/60000 [=====] - 9s 142us/step - loss: 0.1421
 - acc: 0.9579 - val_loss: 0.0873 - val_acc: 0.9732

Epoch 6/10

60000/60000 [=====] - 8s 134us/step - loss: 0.1291
 - acc: 0.9620 - val_loss: 0.0882 - val_acc: 0.9746

Epoch 7/10

60000/60000 [=====] - 8s 134us/step - loss: 0.1187
 - acc: 0.9644 - val_loss: 0.0754 - val_acc: 0.9783

Epoch 8/10

60000/60000 [=====] - 8s 137us/step - loss: 0.1099
 - acc: 0.9679 - val_loss: 0.0712 - val_acc: 0.9790

Epoch 9/10

60000/60000 [=====] - 8s 137us/step - loss: 0.1017
 - acc: 0.9699 - val_loss: 0.0657 - val_acc: 0.9809

Epoch 10/10

60000/60000 [=====] - 8s 140us/step - loss: 0.0973
 - acc: 0.9709 - val_loss: 0.0675 - val_acc: 0.9790

In [14]:


```

1 score = model_2.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06748124550754438

Test accuracy: 0.979



In [16]:


```

1  h1_w = w_after[0].flatten().reshape(-1,1)
2  h2_w = w_after[6].flatten().reshape(-1,1)
3  h3_w = w_after[12].flatten().reshape(-1,1)
4  out_w = w_after[18].flatten().reshape(-1,1)
5
6
7  fig = plt.figure()
8  plt.title("Weight matrices after model trained")
9  plt.subplot(1, 4, 1)
10 plt.title("Trained model Weights")
11 ax = sns.violinplot(y=h1_w,color='b')
12 plt.xlabel('Hidden Layer 1')
13
14 plt.subplot(1, 4, 2)
15 plt.title("Trained model Weights")
16 ax = sns.violinplot(y=h2_w, color='r')
17 plt.xlabel('Hidden Layer 2 ')
18
19 plt.subplot(1, 4, 3)
20 plt.title("Trained model Weights")
21 ax = sns.violinplot(y=h3_w, color='r')
22 plt.xlabel('Hidden Layer 3')
23
24 plt.subplot(1, 4, 4)
25 plt.title("Trained model Weights")
26 ax = sns.violinplot(y=out_w,color='y')
27 plt.xlabel('Output Layer ')
28 plt.show()
29
30 del(model_2)

```

C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\scipy\stats\stat
s.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional in
dexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the fu
ture this will be interpreted as an array index, `arr[np.array(seq)]`, which
will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



7.2.3 MLP (25 hidden layers) + Batch-Norm + Dropout + AdamOptimizer

In [17]:

```
1 model_3 = Sequential()
2
3
4 model_3.add(Dense(200, activation='relu', input_shape=(input_dim,), kernel_initializer=
5 model_3.add(BatchNormalization())
6 model_3.add(Dropout(0.5))
7
8 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
9 model_3.add(BatchNormalization())
10 model_3.add(Dropout(0.5))
11
12 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
13 model_3.add(BatchNormalization())
14 model_3.add(Dropout(0.5))
15
16 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
17 model_3.add(BatchNormalization())
18 model_3.add(Dropout(0.5))
19
20 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
21 model_3.add(BatchNormalization())
22 model_3.add(Dropout(0.5))
23
24 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
25 model_3.add(BatchNormalization())
26 model_3.add(Dropout(0.5))
27
28 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
29 model_3.add(BatchNormalization())
30 model_3.add(Dropout(0.5))
31
32 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
33 model_3.add(BatchNormalization())
34 model_3.add(Dropout(0.5))
35
36 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
37 model_3.add(BatchNormalization())
38 model_3.add(Dropout(0.5))
39
40 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
41 model_3.add(BatchNormalization())
42 model_3.add(Dropout(0.5))
43
44 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
45 model_3.add(BatchNormalization())
46 model_3.add(Dropout(0.5))
47
48 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
49 model_3.add(BatchNormalization())
50 model_3.add(Dropout(0.5))
51
52 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
53 model_3.add(BatchNormalization())
54 model_3.add(Dropout(0.5))
55
56 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
57 model_3.add(BatchNormalization())
```

```

58 model_3.add(Dropout(0.5))
59
60 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
61 model_3.add(BatchNormalization())
62 model_3.add(Dropout(0.5))
63
64 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
65 model_3.add(BatchNormalization())
66 model_3.add(Dropout(0.5))
67
68 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
69 model_3.add(BatchNormalization())
70 model_3.add(Dropout(0.5))
71
72 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
73 model_3.add(BatchNormalization())
74 model_3.add(Dropout(0.5))
75
76 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
77 model_3.add(BatchNormalization())
78 model_3.add(Dropout(0.5))
79
80 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
81 model_3.add(BatchNormalization())
82 model_3.add(Dropout(0.5))
83
84 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
85 model_3.add(BatchNormalization())
86 model_3.add(Dropout(0.5))
87
88 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
89 model_3.add(BatchNormalization())
90 model_3.add(Dropout(0.5))
91
92 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
93 model_3.add(BatchNormalization())
94 model_3.add(Dropout(0.5))
95
96 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
97 model_3.add(BatchNormalization())
98 model_3.add(Dropout(0.5))
99
100 model_3.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
101 model_3.add(BatchNormalization())
102 model_3.add(Dropout(0.5))
103
104 model_3.add(Dense(output_dim, activation='softmax'))
105
106 model_3.summary()
107
108

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 200)	157000
batch_normalization_6 (Batch Normalization)	(None, 200)	800
dropout_6 (Dropout)	(None, 200)	0

dense_9 (Dense)	(None, 100)	20100
batch_normalization_7 (Batch Normalization)	(None, 100)	400
dropout_7 (Dropout)	(None, 100)	0
dense_10 (Dense)	(None, 100)	10100
batch_normalization_8 (Batch Normalization)	(None, 100)	400

In [18]:

```
1 model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
2
3 history = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 29s 477us/step - loss: 2.6718
- acc: 0.0996 - val_loss: 2.3042 - val_acc: 0.1135

Epoch 2/10

60000/60000 [=====] - 19s 318us/step - loss: 2.3356
- acc: 0.1038 - val_loss: 2.3017 - val_acc: 0.1135

Epoch 3/10

60000/60000 [=====] - 19s 321us/step - loss: 2.3108
- acc: 0.1067 - val_loss: 2.3025 - val_acc: 0.1135

Epoch 4/10

60000/60000 [=====] - 19s 321us/step - loss: 2.3074
- acc: 0.1067 - val_loss: 2.3019 - val_acc: 0.1135

Epoch 5/10

60000/60000 [=====] - 19s 320us/step - loss: 2.3071
- acc: 0.1086 - val_loss: 2.3027 - val_acc: 0.1135

Epoch 6/10

60000/60000 [=====] - 19s 320us/step - loss: 2.3073
- acc: 0.1057 - val_loss: 2.3031 - val_acc: 0.1135

Epoch 7/10

60000/60000 [=====] - 19s 322us/step - loss: 2.3069
- acc: 0.1084 - val_loss: 2.3018 - val_acc: 0.1135

Epoch 8/10

60000/60000 [=====] - 19s 324us/step - loss: 2.3055
- acc: 0.1082 - val_loss: 2.3021 - val_acc: 0.1135

Epoch 9/10

60000/60000 [=====] - 22s 369us/step - loss: 2.3064
- acc: 0.1053 - val_loss: 2.3022 - val_acc: 0.1135

Epoch 10/10

60000/60000 [=====] - 20s 337us/step - loss: 2.3058
- acc: 0.1056 - val_loss: 2.3020 - val_acc: 0.1135

In [19]:

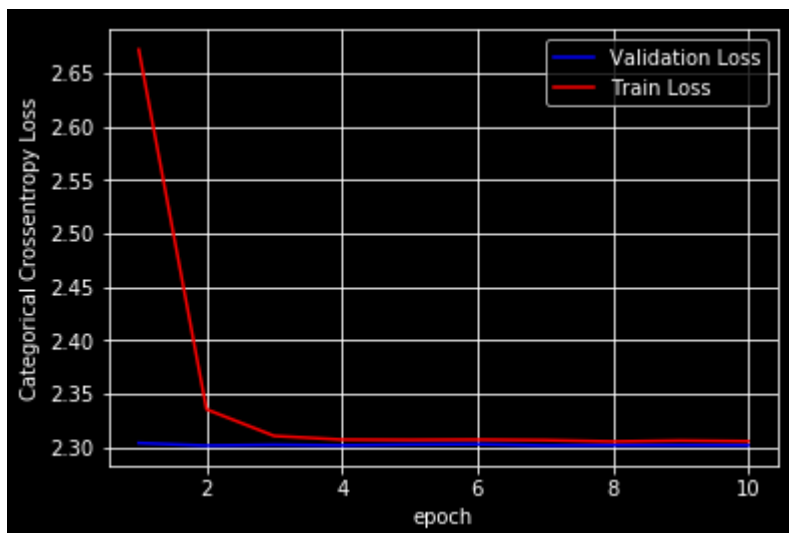
```

1 score = model_3.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)
26
27 del(model_3)

```

Test score: 2.302019982147217

Test accuracy: 0.1135



7.2.4 MLP (4 hidden layers) + Batch-Norm + Dropout + AdamOptimizer

In [20]:

```

1 model_4 = Sequential()
2
3 model_4.add(Dense(100, activation='relu', input_shape=(input_dim,), kernel_initializer='he_normal'))
4 model_4.add(BatchNormalization())
5 model_4.add(Dropout(0.5))
6
7 model_4.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
8 model_4.add(BatchNormalization())
9 model_4.add(Dropout(0.5))
10
11 model_4.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
12 model_4.add(BatchNormalization())
13 model_4.add(Dropout(0.5))
14
15 model_4.add(Dense(100, activation='relu', kernel_initializer='he_normal'))
16 model_4.add(BatchNormalization())
17 model_4.add(Dropout(0.5))
18
19 model_4.add(Dense(output_dim, activation='softmax'))
20
21 model_4.summary()

```

Layer (type)	Output Shape	Param #
=====		
dense_34 (Dense)	(None, 100)	78500
batch_normalization_31 (Batch Normalization)	(None, 100)	400
dropout_31 (Dropout)	(None, 100)	0
dense_35 (Dense)	(None, 100)	10100
batch_normalization_32 (Batch Normalization)	(None, 100)	400
dropout_32 (Dropout)	(None, 100)	0
dense_36 (Dense)	(None, 100)	10100
batch_normalization_33 (Batch Normalization)	(None, 100)	400
dropout_33 (Dropout)	(None, 100)	0
dense_37 (Dense)	(None, 100)	10100
batch_normalization_34 (Batch Normalization)	(None, 100)	400
dropout_34 (Dropout)	(None, 100)	0
dense_38 (Dense)	(None, 10)	1010
=====		
Total params: 111,410		
Trainable params: 110,610		
Non-trainable params: 800		

In [21]:

```

1 model_4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 8s 132us/step - loss: 1.3473

- acc: 0.5643 - val_loss: 0.3420 - val_acc: 0.9041

Epoch 2/10

60000/60000 [=====] - 5s 81us/step - loss: 0.5937 -

acc: 0.8200 - val_loss: 0.2426 - val_acc: 0.9299

Epoch 3/10

60000/60000 [=====] - 5s 82us/step - loss: 0.4655 -

acc: 0.8671 - val_loss: 0.2026 - val_acc: 0.9411

Epoch 4/10

60000/60000 [=====] - 5s 77us/step - loss: 0.3861 -

acc: 0.8914 - val_loss: 0.1887 - val_acc: 0.9482

Epoch 5/10

60000/60000 [=====] - 5s 77us/step - loss: 0.3411 -

acc: 0.9063 - val_loss: 0.1684 - val_acc: 0.9521

Epoch 6/10

60000/60000 [=====] - 5s 76us/step - loss: 0.3097 -

acc: 0.9158 - val_loss: 0.1519 - val_acc: 0.9575

Epoch 7/10

60000/60000 [=====] - 5s 78us/step - loss: 0.2953 -

acc: 0.9196 - val_loss: 0.1446 - val_acc: 0.9595

Epoch 8/10

60000/60000 [=====] - 5s 80us/step - loss: 0.2778 -

acc: 0.9250 - val_loss: 0.1413 - val_acc: 0.9611

Epoch 9/10

60000/60000 [=====] - 5s 78us/step - loss: 0.2596 -

acc: 0.9294 - val_loss: 0.1308 - val_acc: 0.9654

Epoch 10/10

60000/60000 [=====] - 5s 79us/step - loss: 0.2502 -

acc: 0.9319 - val_loss: 0.1341 - val_acc: 0.9650

In [22]:

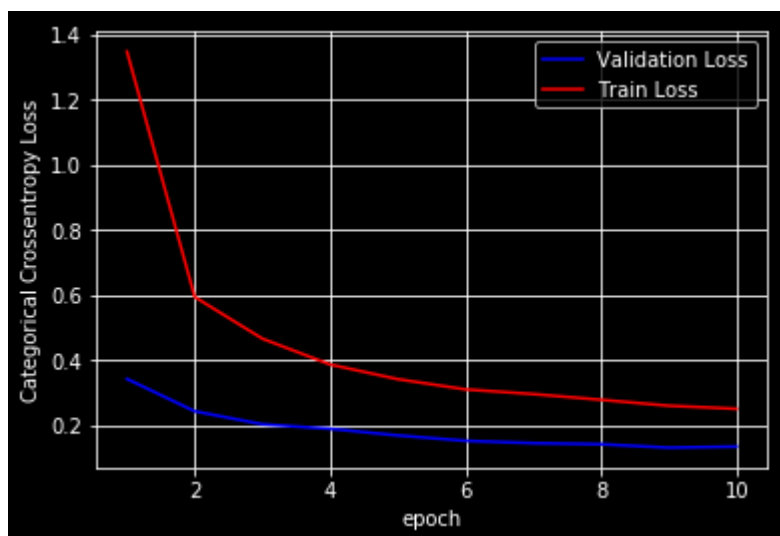
```

1 score = model_4.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.13409332928974182

Test accuracy: 0.965



In [24]:

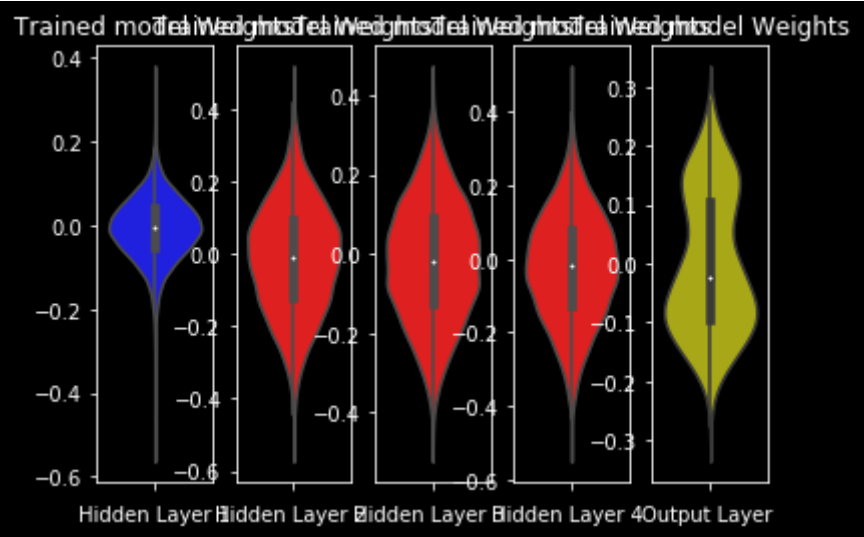
```

1  h1_w = w_after[0].flatten().reshape(-1,1)
2  h2_w = w_after[6].flatten().reshape(-1,1)
3  h3_w = w_after[12].flatten().reshape(-1,1)
4  h4_w = w_after[18].flatten().reshape(-1,1)
5  out_w = w_after[24].flatten().reshape(-1,1)
6
7
8  fig = plt.figure()
9  plt.title("Weight matrices after model trained")
10 plt.subplot(1, 5, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 5, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 5, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=h3_w, color='r')
23 plt.xlabel('Hidden Layer 3 ')
24
25 plt.subplot(1, 5, 4)
26 plt.title("Trained model Weights")
27 ax = sns.violinplot(y=h4_w, color='r')
28 plt.xlabel('Hidden Layer 4 ')
29
30 plt.subplot(1, 5, 5)
31 plt.title("Trained model Weights")
32 ax = sns.violinplot(y=out_w,color='y')
33 plt.xlabel('Output Layer ')
34 plt.show()
35
36 del(model_4)

```

C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



7.2.5 MLP (1 hidden layer) + Batch-Norm + Dropout + AdamOptimizer

In [25]:

```
1 model_5 = Sequential()
2
3 model_5.add(Dense(1000, activation='relu', input_shape=(input_dim,), kernel_initializer=
4 model_5.add(BatchNormalization())
5 model_5.add(Dropout(0.5))
6
7 model_5.add(Dense(output_dim, activation='softmax'))
8
9 model_5.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_39 (Dense)	(None, 1000)	785000
batch_normalization_35 (Batch Normalization)	(None, 1000)	4000
dropout_35 (Dropout)	(None, 1000)	0
dense_40 (Dense)	(None, 10)	10010
=====	=====	=====
Total params: 799,010		
Trainable params: 797,010		
Non-trainable params: 2,000		

In [26]:

```

1 model_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 12s 206us/step - loss: 0.2689

- acc: 0.9215 - val_loss: 0.1091 - val_acc: 0.9662

Epoch 2/10

60000/60000 [=====] - 10s 164us/step - loss: 0.1248

- acc: 0.9616 - val_loss: 0.0872 - val_acc: 0.9743

Epoch 3/10

60000/60000 [=====] - 10s 165us/step - loss: 0.0962

- acc: 0.9700 - val_loss: 0.0806 - val_acc: 0.9748

Epoch 4/10

60000/60000 [=====] - 10s 166us/step - loss: 0.0795

- acc: 0.9747 - val_loss: 0.0689 - val_acc: 0.9783

Epoch 5/10

60000/60000 [=====] - 10s 169us/step - loss: 0.0725

- acc: 0.9763 - val_loss: 0.0653 - val_acc: 0.9806

Epoch 6/10

60000/60000 [=====] - 10s 167us/step - loss: 0.0629

- acc: 0.9800 - val_loss: 0.0673 - val_acc: 0.9797

Epoch 7/10

60000/60000 [=====] - 10s 172us/step - loss: 0.0610

- acc: 0.9798 - val_loss: 0.0622 - val_acc: 0.9808

Epoch 8/10

60000/60000 [=====] - 10s 172us/step - loss: 0.0567

- acc: 0.9806 - val_loss: 0.0637 - val_acc: 0.9819

Epoch 9/10

60000/60000 [=====] - 10s 175us/step - loss: 0.0503

- acc: 0.9831 - val_loss: 0.0647 - val_acc: 0.9816

Epoch 10/10

60000/60000 [=====] - 10s 175us/step - loss: 0.0496

- acc: 0.9836 - val_loss: 0.0579 - val_acc: 0.9822

In [27]:


```

1 score = model_5.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05786827448041295

Test accuracy: 0.9822



In [29]:


```

1 w_after = model_5.get_weights()
2 h1_w = w_after[0].flatten().reshape(-1,1)
3 out_w = w_after[6].flatten().reshape(-1,1)
4
5
6 fig = plt.figure()
7 plt.title("Weight matrices after model trained")
8 plt.subplot(1, 3, 1)
9 plt.title("Trained model Weights")
10 ax = sns.violinplot(y=h1_w,color='b')
11 plt.xlabel('Hidden Layer 1')
12
13 plt.subplot(1, 3, 3)
14 plt.title("Trained model Weights")
15 ax = sns.violinplot(y=out_w,color='y')
16 plt.xlabel('Output Layer ')
17 plt.show()
18
19 del(model_5)

```

C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



7.2.6 MLP (1 hidden layer) + Batch-Norm + Dropout (low) + AdamOptimizer

In [30]:

```
1 model_6 = Sequential()
2
3 model_6.add(Dense(300, activation='relu', input_shape=(input_dim,), kernel_initializer=
4 model_6.add(BatchNormalization())
5 model_6.add(Dropout(0.1))
6
7 model_6.add(Dense(output_dim, activation='softmax'))
8
9 model_6.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_41 (Dense)	(None, 300)	235500
=====		
batch_normalization_36 (Batch Normalization)	(None, 300)	1200
=====		
dropout_36 (Dropout)	(None, 300)	0
=====		
dense_42 (Dense)	(None, 10)	3010
=====		
Total params: 239,710		
Trainable params: 239,110		
Non-trainable params: 600		
=====		

In [31]:

```

1 model_6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 98us/step - loss: 0.2407 -
acc: 0.9280 - val_loss: 0.1131 - val_acc: 0.9678

Epoch 2/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0964 -
acc: 0.9715 - val_loss: 0.0853 - val_acc: 0.9743

Epoch 3/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0629 -
acc: 0.9813 - val_loss: 0.0797 - val_acc: 0.9759

Epoch 4/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0474 -
acc: 0.9863 - val_loss: 0.0749 - val_acc: 0.9767

Epoch 5/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0396 -
acc: 0.9877 - val_loss: 0.0695 - val_acc: 0.9793

Epoch 6/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0334 -
acc: 0.9897 - val_loss: 0.0798 - val_acc: 0.9771

Epoch 7/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0255 -
acc: 0.9922 - val_loss: 0.0719 - val_acc: 0.9782

Epoch 8/10

60000/60000 [=====] - 4s 67us/step - loss: 0.0217 -
acc: 0.9933 - val_loss: 0.0746 - val_acc: 0.9768

Epoch 9/10

60000/60000 [=====] - 4s 66us/step - loss: 0.0207 -
acc: 0.9934 - val_loss: 0.0754 - val_acc: 0.9790

Epoch 10/10

60000/60000 [=====] - 4s 67us/step - loss: 0.0206 -
acc: 0.9934 - val_loss: 0.0693 - val_acc: 0.9800

In [32]:


```

1 score = model_6.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06930483953608782

Test accuracy: 0.98



In [34]:

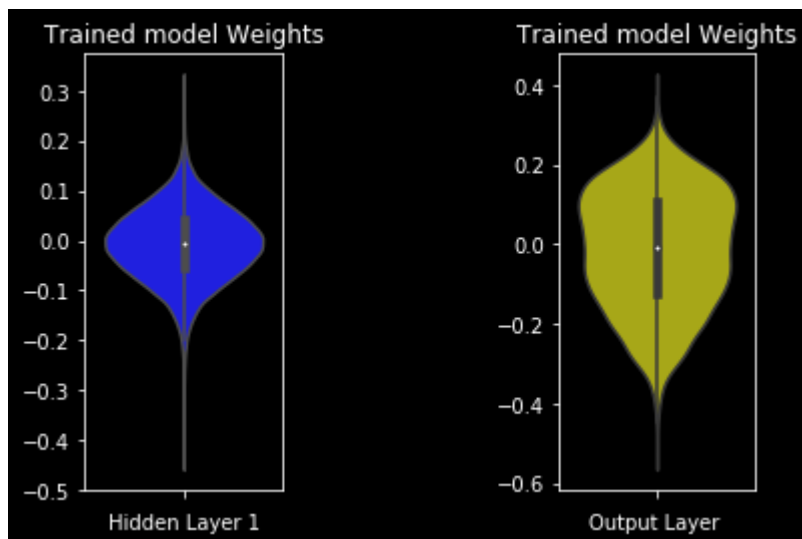
```

1 w_after = model_6.get_weights()
2 h1_w = w_after[0].flatten().reshape(-1,1)
3 out_w = w_after[6].flatten().reshape(-1,1)
4
5
6 fig = plt.figure()
7 plt.title("Weight matrices after model trained")
8 plt.subplot(1, 3, 1)
9 plt.title("Trained model Weights")
10 ax = sns.violinplot(y=h1_w,color='b')
11 plt.xlabel('Hidden Layer 1')
12
13 plt.subplot(1, 3, 3)
14 plt.title("Trained model Weights")
15 ax = sns.violinplot(y=out_w,color='y')
16 plt.xlabel('Output Layer ')
17 plt.show()
18
19 del(model_6)

```

C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



7.2.7 MLP (1 hidden layer) + Batch-Norm + Dropout (high) + AdamOptimizer

In [35]:

```
1 model_7 = Sequential()
2
3 model_7.add(Dense(300, activation='relu', input_shape=(input_dim,), kernel_initializer=
4 model_7.add(BatchNormalization())
5 model_7.add(Dropout(0.9))
6
7 model_7.add(Dense(output_dim, activation='softmax'))
8
9 model_7.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_43 (Dense)	(None, 300)	235500
batch_normalization_37 (Batch Normalization)	(None, 300)	1200
dropout_37 (Dropout)	(None, 300)	0
dense_44 (Dense)	(None, 10)	3010
=====		
Total params: 239,710		
Trainable params: 239,110		
Non-trainable params: 600		

In [36]:

```

1 model_7.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
2
3 history = model_7.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 104us/step - loss: 1.0630
- acc: 0.7194 - val_loss: 0.2704 - val_acc: 0.9217

Epoch 2/10

60000/60000 [=====] - 4s 68us/step - loss: 0.5531 -
acc: 0.8324 - val_loss: 0.2353 - val_acc: 0.9301

Epoch 3/10

60000/60000 [=====] - 4s 68us/step - loss: 0.4699 -
acc: 0.8575 - val_loss: 0.2184 - val_acc: 0.9344

Epoch 4/10

60000/60000 [=====] - 4s 68us/step - loss: 0.4361 -
acc: 0.8693 - val_loss: 0.2031 - val_acc: 0.9395

Epoch 5/10

60000/60000 [=====] - 4s 69us/step - loss: 0.4146 -
acc: 0.8765 - val_loss: 0.1903 - val_acc: 0.9399

Epoch 6/10

60000/60000 [=====] - 4s 69us/step - loss: 0.3932 -
acc: 0.8820 - val_loss: 0.1848 - val_acc: 0.9442

Epoch 7/10

60000/60000 [=====] - 4s 69us/step - loss: 0.3817 -
acc: 0.8842 - val_loss: 0.1725 - val_acc: 0.9470

Epoch 8/10

60000/60000 [=====] - 4s 70us/step - loss: 0.3709 -
acc: 0.8880 - val_loss: 0.1705 - val_acc: 0.9483

Epoch 9/10

60000/60000 [=====] - 4s 69us/step - loss: 0.3637 -
acc: 0.8915 - val_loss: 0.1619 - val_acc: 0.9510

Epoch 10/10

60000/60000 [=====] - 4s 69us/step - loss: 0.3562 -
acc: 0.8944 - val_loss: 0.1598 - val_acc: 0.9525

In [37]:


```

1 score = model_7.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # List of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_data
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to number of ep
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.15979730788581073

Test accuracy: 0.9525



In [39]:


```

1 w_after = model_7.get_weights()
2 h1_w = w_after[0].flatten().reshape(-1,1)
3 out_w = w_after[6].flatten().reshape(-1,1)
4
5
6 fig = plt.figure()
7 plt.title("Weight matrices after model trained")
8 plt.subplot(1, 3, 1)
9 plt.title("Trained model Weights")
10 ax = sns.violinplot(y=h1_w,color='b')
11 plt.xlabel('Hidden Layer 1')
12
13 plt.subplot(1, 3, 3)
14 plt.title("Trained model Weights")
15 ax = sns.violinplot(y=out_w,color='y')
16 plt.xlabel('Output Layer ')
17 plt.show()
18
19 del(model_7)

```

C:\Users\Byron\AppData\Local\PythonMaster\lib\site-packages\scipy\stats\stat
s.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional in
dexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the fu
ture this will be interpreted as an array index, `arr[np.array(seq)]`, which
will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



Type Markdown and LaTeX: α^2

In [2]:

```

1 import prettytable
2
3 table = prettytable.PrettyTable()
4 table.add_row(['Model', 'Parameters', 'Training loss', 'Validation loss', 'Accuracy', 'epochs'])
5 table.add_row(['2 HL,relu,adam,dropout,batch norm', '471434', '0.08', '0.06', '98', '10'])
6 table.add_row(['3 HL,relu,adam,dropout,batch norm', '577520', '0.09', '0.07', '97', '10'])
7 table.add_row(['25 HL,relu,adam,dropout,batch norm,100 neurons', '420810', '2.3', '2.3', '11', '10'])
8 table.add_row(['4 HL,relu,adam,dropout,batch norm,100 neurons', '111410', '0.17', '0.09', '97', '10'])
9 table.add_row(['1 HL,relu,adam,dropout,batch norm,1000 neurons', '799010', '0.05', '0.06', '98', '10'])
10 table.add_row(['1 HL,relu,adam,batch norm,low dropout,300 neurons', '239710', '0.01', '0.06', '98', '10'])
11 table.add_row(['1 HL,relu,adam,batch norm,high dropout,300 neurons', '239710', '0.36', '0.16', '95', '10'])
12
13 print(table)

```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|               | Field 1 |               | Field 2 |               | Field 3 |
|               | Field 4 | Field 5 | Field 6 |               | Field 7 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|               | Model |               | Parameters | Training |
|               | loss | Validation loss | Accuracy | epochs |
|               | 2 HL,relu,adam,dropout,batch norm | 471434 | 0.0 |
8 |               | 0.06 | 98 | 10 |
|               | 3 HL,relu,adam,dropout,batch norm | 577520 | 0.0 |
9 |               | 0.07 | 97 | 10 |
|               | 25 HL,relu,adam,dropout,batch norm,100 neurons | 420810 | 2.3 |
|               | 2.3 | 11 | 10 |
|               | 4 HL,relu,adam,dropout,batch norm,100 neurons | 111410 | 0.1 |
7 |               | 0.09 | 97 | 10 |
|               | 1 HL,relu,adam,dropout,batch norm,1000 neurons | 799010 | 0.0 |
5 |               | 0.06 | 98 | 10 |
|               | 1 HL,relu,adam,batch norm,low dropout,300 neurons | 239710 | 0.0 |
1 |               | 0.06 | 98 | 10 |
|               | 1 HL,relu,adam,batch norm,high dropout,300 neurons | 239710 | 0.3 |
6 |               | 0.16 | 95 | 10 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+

```